



# Swin Transformer: Hierarchical Vision Transformer using Shifted Windows

- This paper presents a new vision transformer, called Swin Transformer, that capably serves as a **general-purpose backbone** for computer vision.
- Challenges in adapting Transformer from language to vision: **high resolution of pixels in images** compared to **words in text**.
- The proposed Swin Transformer **merge image patches** in deeper layers. **Computational complexity** has **linear** to input **image size** due to computation of self-attention only within **each local window**.

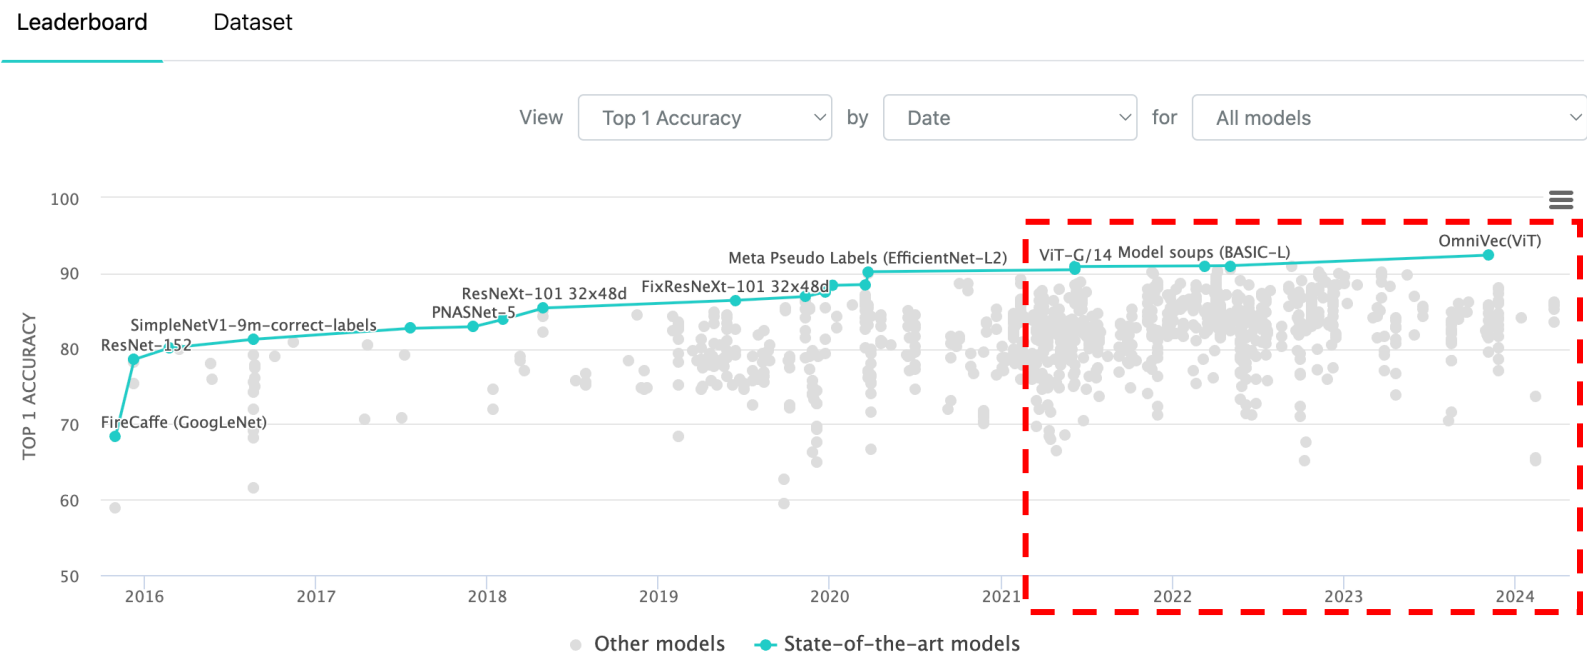


# Swin Transformer: Hierarchical Vision Transformer using Shifted Windows

## Image Classification on ImageNet




CNNs serve as backbone networks for a variety of vision tasks.



Its tremendous success in the language domain has led researchers to investigate its adaptation to computer vision (image classification and joint vision-language modeling)


# Swin Transformer: Hierarchical Vision Transformer using Shifted Windows

- Swin Transformer constructs a hierarchical representation by starting from small-sized patches and gradually merging neighboring patches in deeper Transformer layers.
- Number of patches in each window is fixed, and thus the complexity becomes linear to image size




# Swin Transformer: Hierarchical Vision Transformer using Shifted Windows

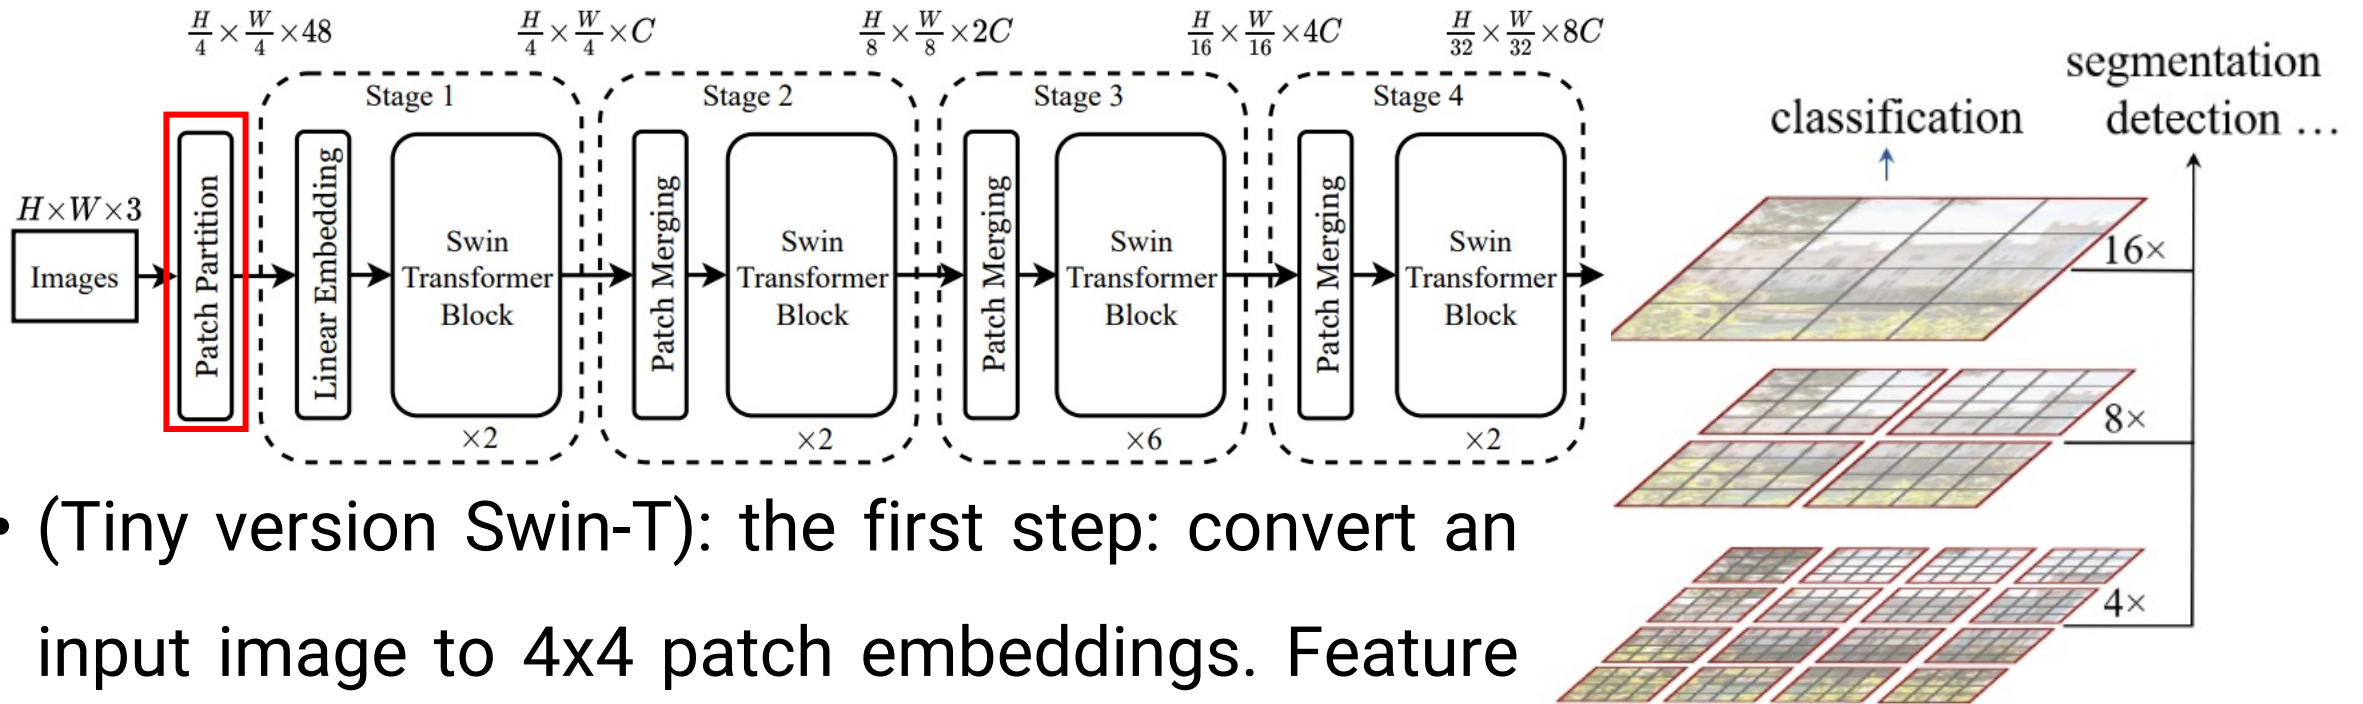
- A key design element of Swin Transformer is its shift of the window partition between consecutive self-attention layers.
- How to explain this figure? (look shifted window - multi self attention)



# Hierarchical feature map




# Swin Transformer: Hierarchical Vision Transformer using Shifted Windows




- (Tiny version Swin-T): the first step: convert an input image to 4x4 patch embeddings. Feature dimension of each patch is  $4 \times 4 \times 3 = 48$

- image size: 224x224
- patch size: 4x4
- number of patches: 56x56
- total patches: 3136
- every patch: 4x4x3 values





# Swin Transformer: Hierarchical Vision Transformer using Shifted Windows




Swin-T:  $C = 96$ , layer numbers = {2, 2, 6, 2}  
Swin-S:  $C = 96$ , layer numbers = {2, 2, 18, 2}  
Swin-B:  $C = 128$ , layer numbers = {2, 2, 18, 2}  
Swin-L:  $C = 192$ , layer numbers = {2, 2, 18, 2}


$C$  is a hyper parameter, usually is 96 or 128

- A linear embedding layer is applied on this raw valued feature to project it to an arbitrary dimension (denoted as  $C$ )

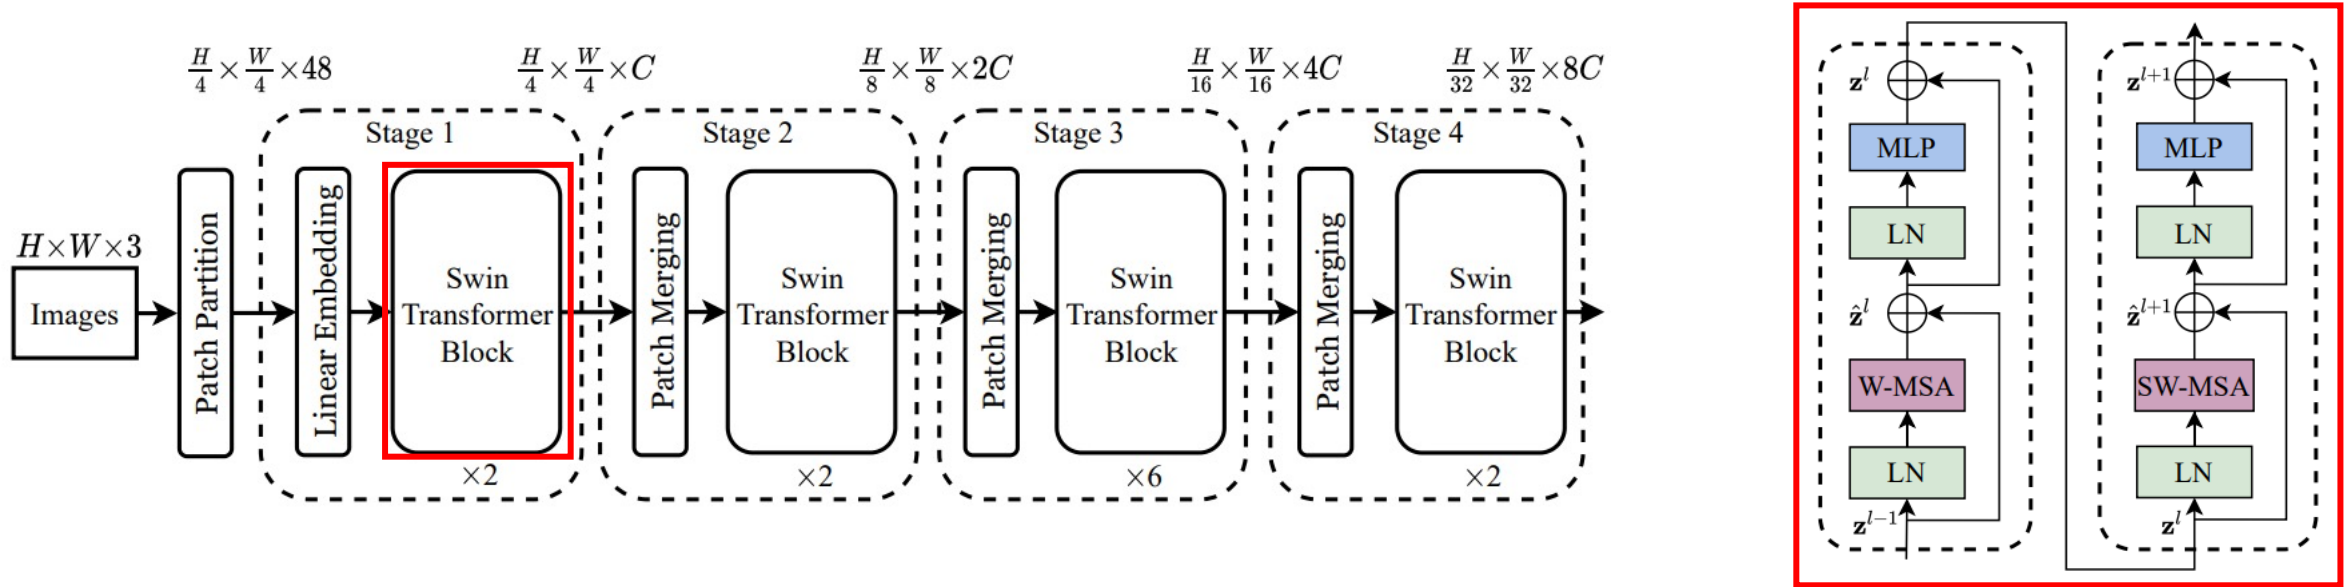
Conv2D (4x4)  
Stride = 4  
Kernel =  $C$



Multi  
Layer  
Perceptron



# Swin Transformer: Hierarchical Vision Transformer using Shifted Windows




- Swin Transformer block consists of a shifted window based **MSA** module, followed by a 2-layer **MLP with GELU** nonlinearity in between. A **LayerNorm** (LN) layer is applied before each MSA module and each MLP, and a **residual connection** is applied after each module.


- W-MSA: Window Multilayer Self Attention
- SW-MSA: Shifted Window Multilayer Self Attention




# Swin Transformer: Hierarchical Vision Transformer using Shifted Windows



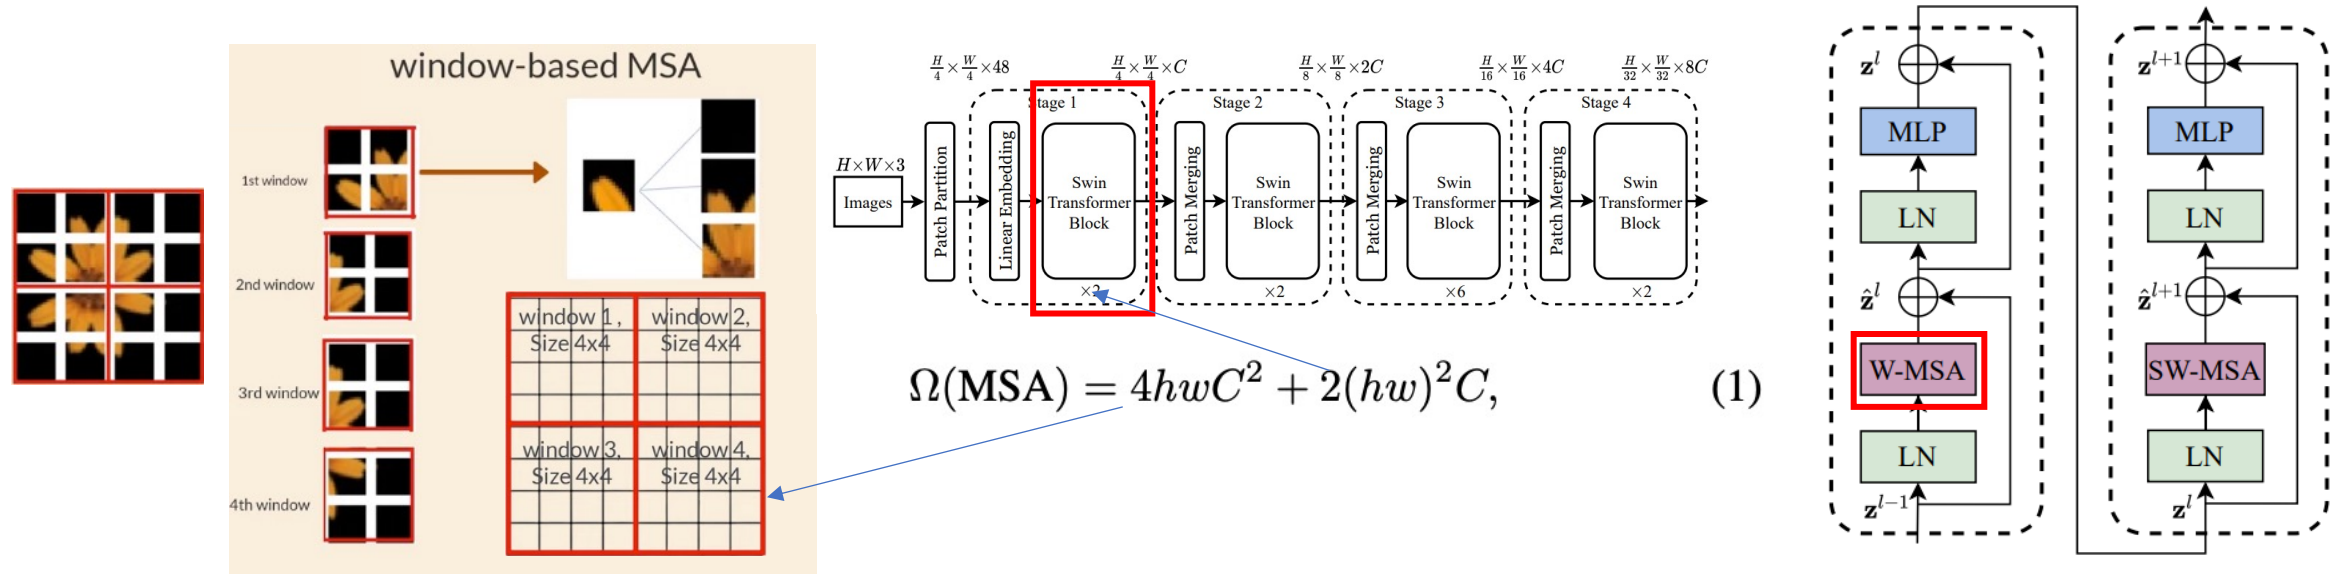
**Attention in ViT:** calculate the relationship between each patch and all the other patches. Computational complexity meaning inefficient in the high resolution images case.  $O(N^2)$ ,  $N$  – number of patches.



- Instead of computing self-attention on all the patches in an input image, window attention is limited to include patches that correspond to a window of predefined shape.




# Swin Transformer: Hierarchical Vision Transformer using Shifted Windows




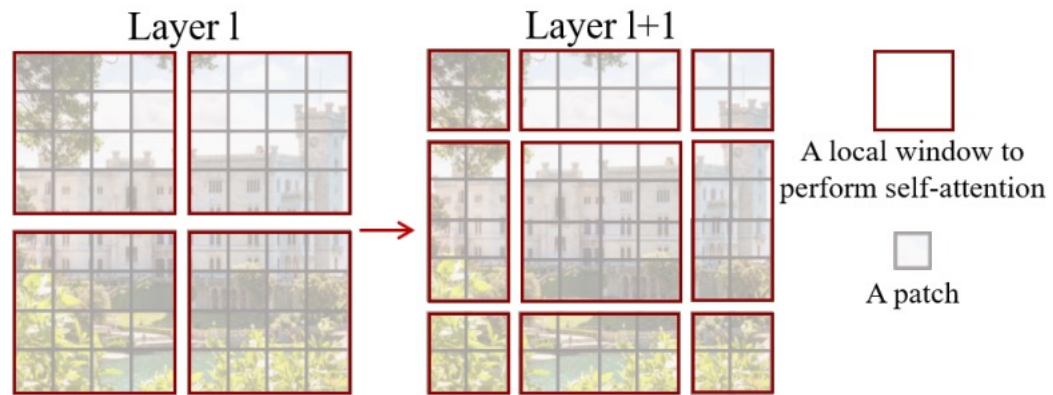

- W-MSA: collection of patches and attention is computed only within each window
- An input image is split into multiple non-overlapping windows of equal size for attention computation.
- The global computation leads to quadratic complexity with respect to the number of tokens.

# Swin Transformer: Hierarchical Vision Transformer using Shifted Windows





- This windows based self-attention cannot provide connections between windows.
- Windows are shifted in consecutive layers by  $M/2$  and  $M/2$  patches horizontally and vertically respectively. In the following figure windows 1, 2, 3 and 4 are shifted by 2 patches as  $M=4$  in this case.

# Swin Transformer: Hierarchical Vision Transformer using Shifted Windows



- shift size =  $M/2$  (M: window size)
- Shift 2 times to the left
- Shift 2 times to the top

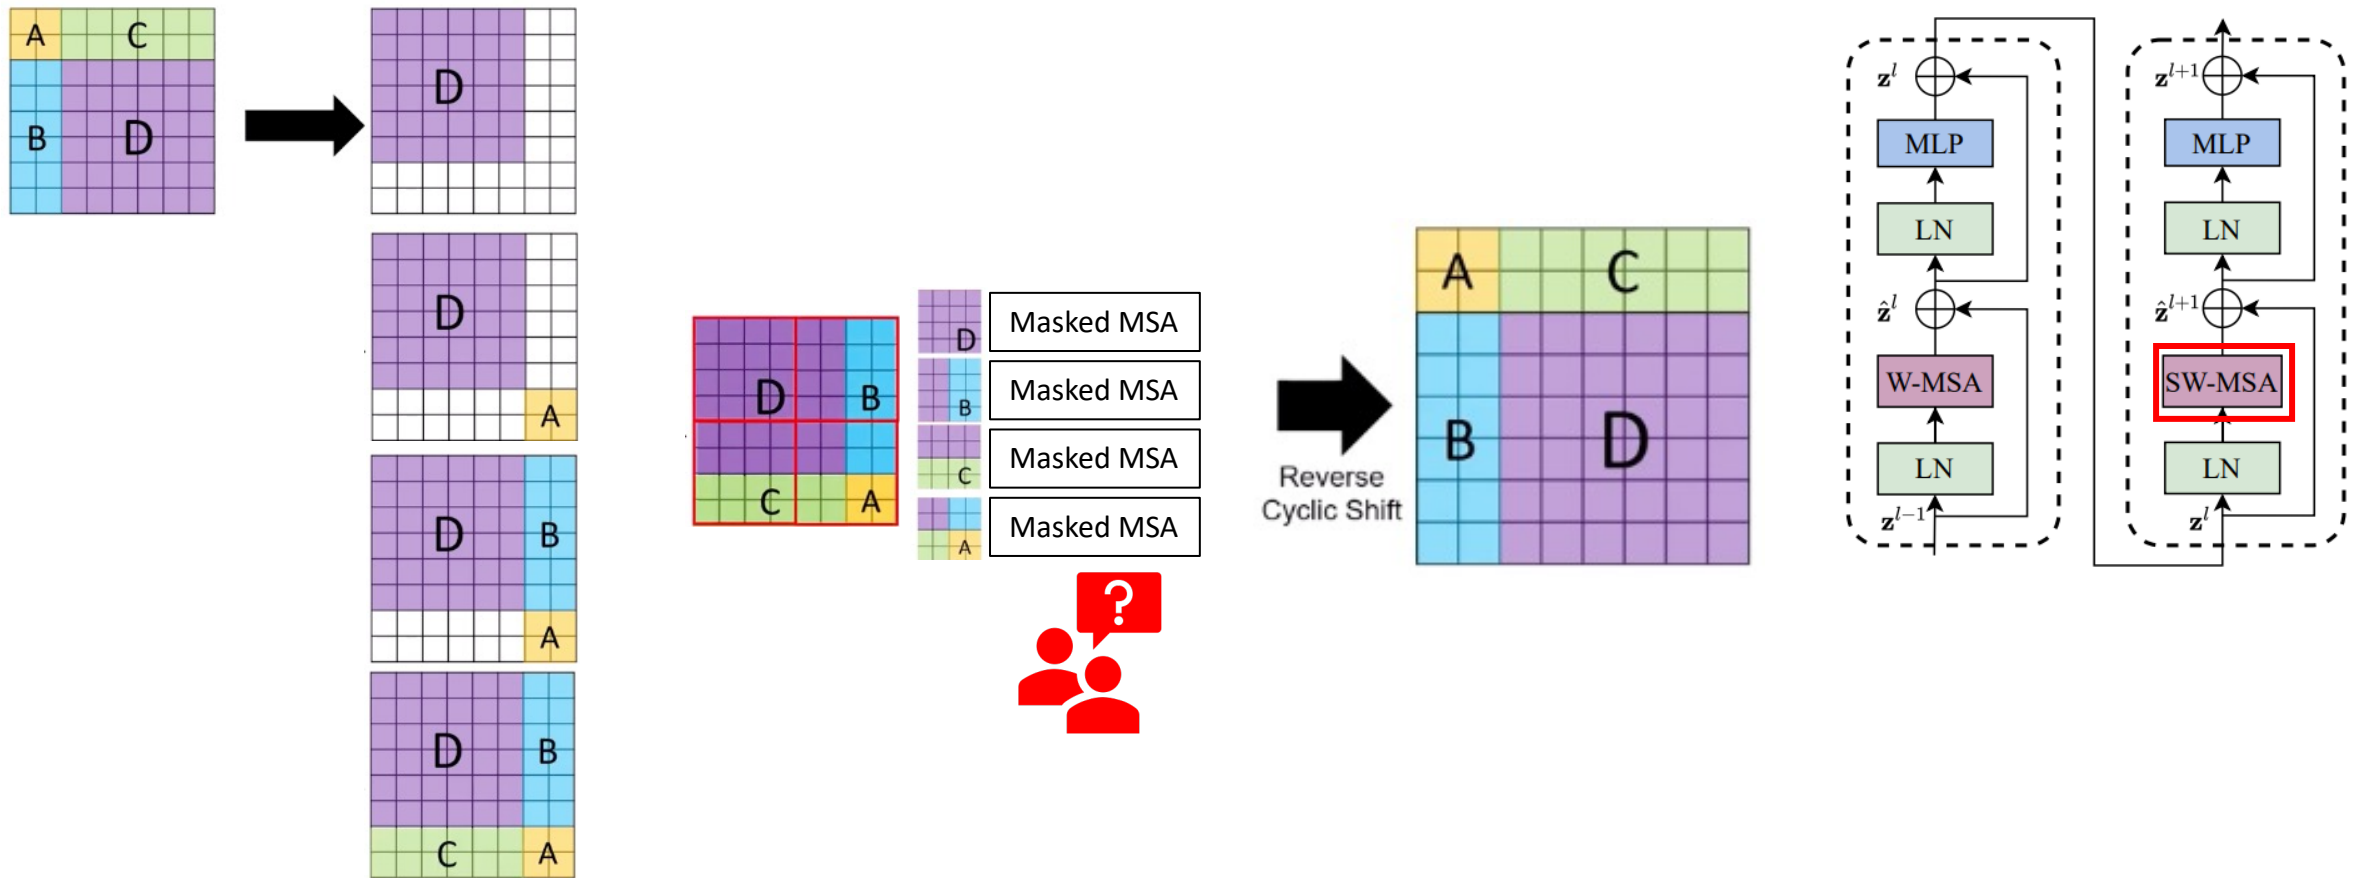
# Swin Transformer: Hierarchical Vision Transformer using Shifted Windows



- **Problems:** zero patch is completely useless to compute relationship in MAS.
- WMSA/SWMSA cannot cover object.



# Swin Transformer: Hierarchical Vision Transformer using Shifted Windows



- **Solution:** cyclic shifting (rolling) the additional windows then perform attention computation on cyclic shifted configuration, and then shifting the windows back to the original locations.

# The Question about the mask of window attention #38



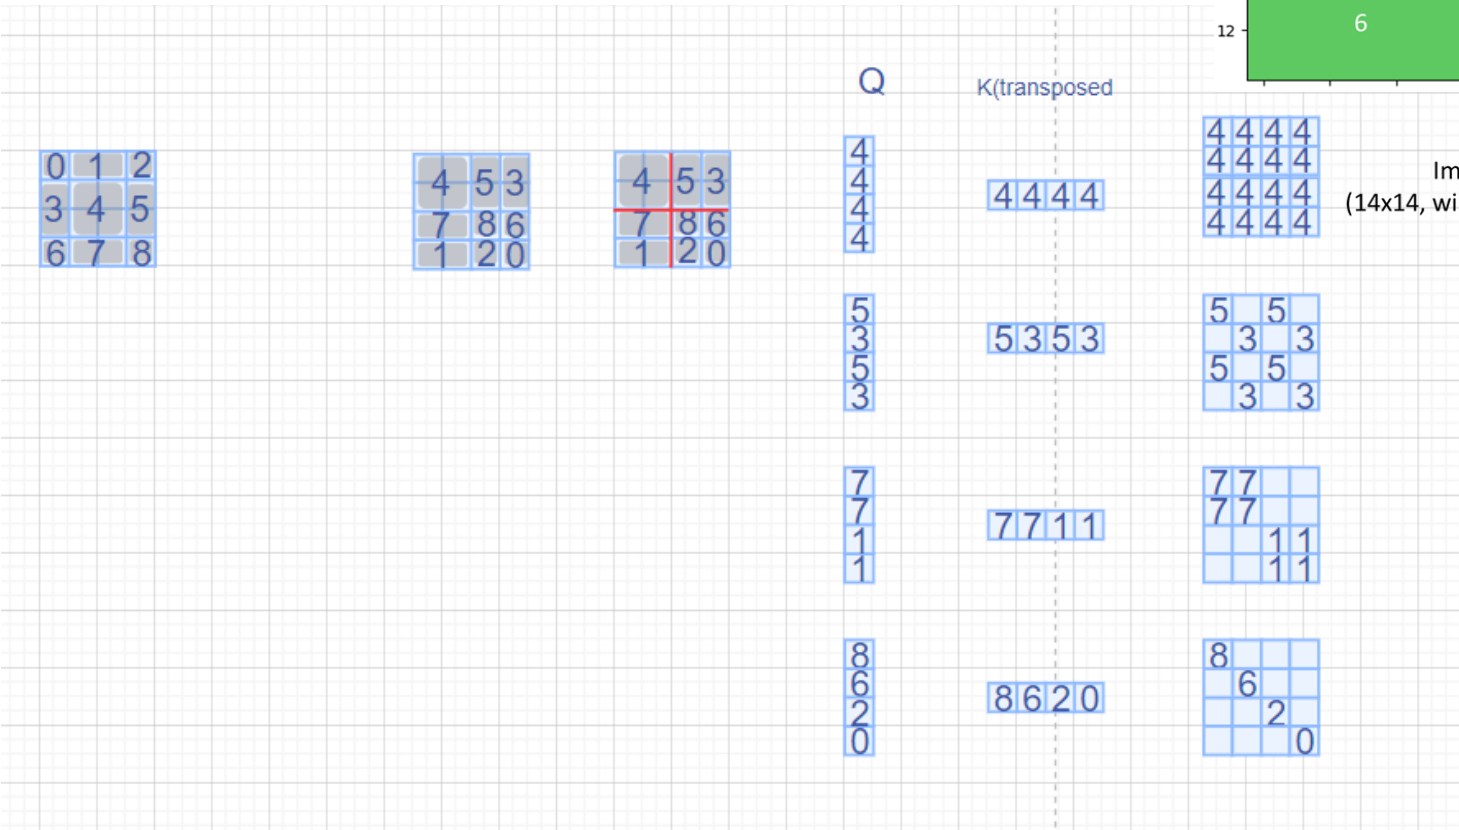



Image Mask  
(14x14, window 7x7, shift 3)

Attn Mask

# Swin Transformer: Hierarchical Vision Transformer using Shifted Windows




## Patch Merging

Assuming that  $n=2$ , and each group consists of  $2 \times 2$  neighboring patches

Step 1: Split input image into groups of  $2 \times 2$

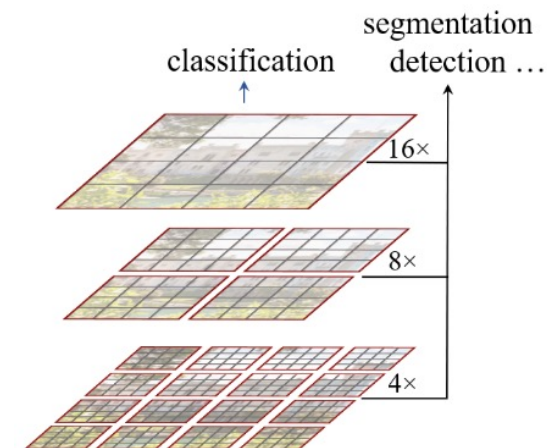

Step 2: In each group, stack the patches depth-wise

Step 3: Combine the stacked groups




- The patch-merging layer merges four patches at a time, so with each merge the height and width of the image are reduced by a factor of 2.
- in stage 1 the input resolution is  $(H/4, W/4)$ , but after patch merging the resolution becomes  $(H/8, W/8)$ , which is the input for stage 2, for stage 3 is  $(H/16, W/16)$  and for stage 4 it is  $(H/32, W/32)$ .

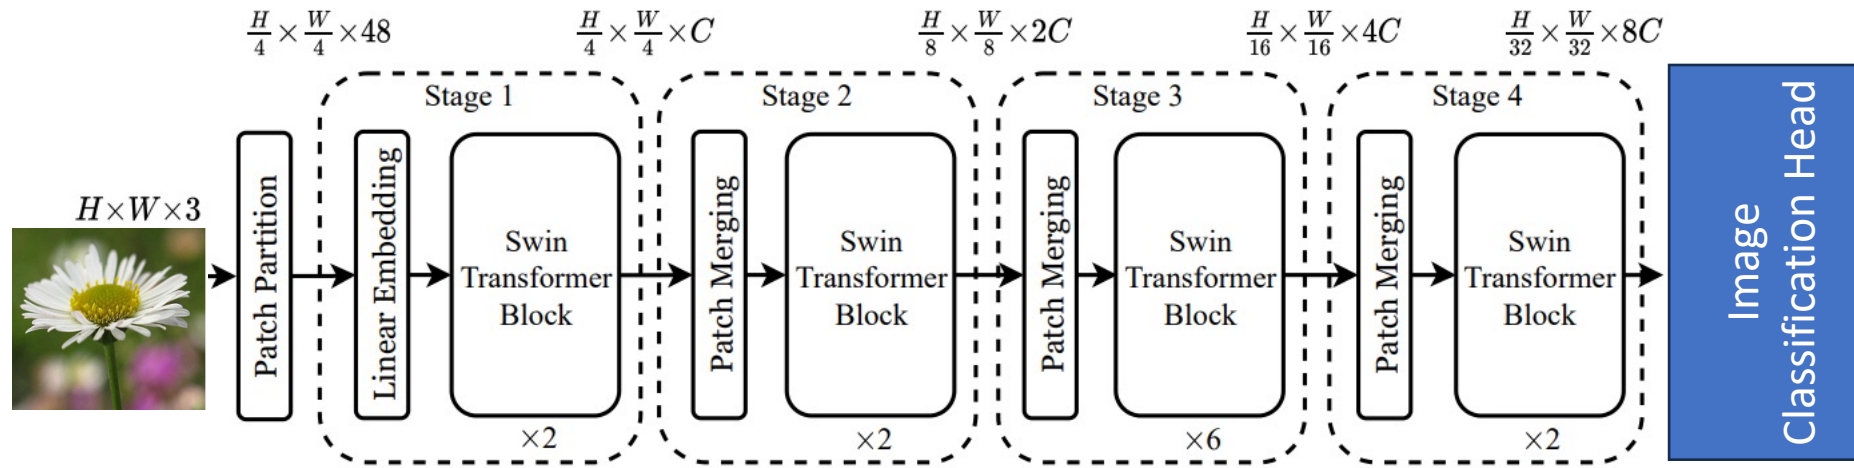
# Swin Transformer: Hierarchical Vision Transformer using Shifted Windows



Split image  
into 2x2 grid

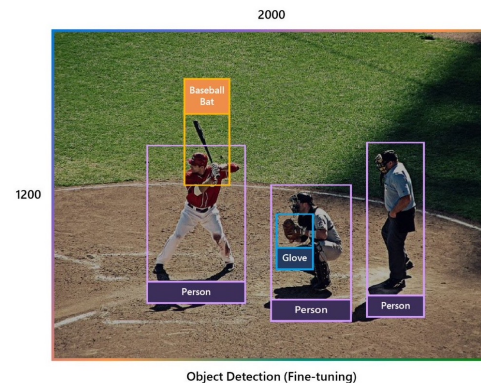
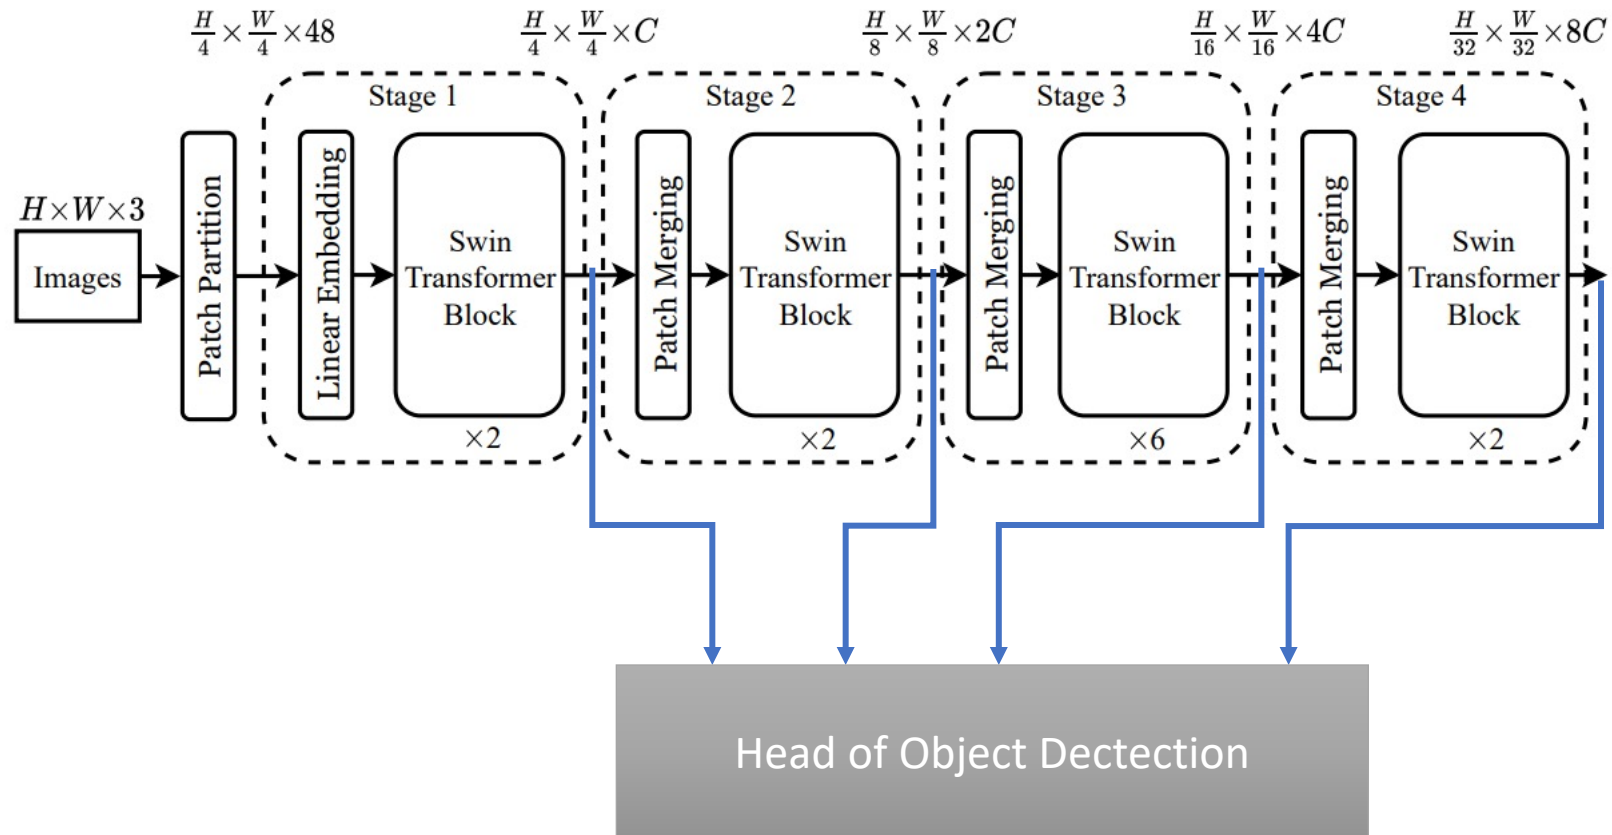


# Swin Transformer: Hierarchical Vision Transformer using Shifted Windows







# Swin Transformer: Hierarchical Vision Transformer using Shifted Windows



- [Swin Transformer: Hierarchical Vision Transformer Using Shifted Windows \(thecvf.com\)](#)
- [A Summary of the Swin Transformer: A Hierarchical Vision Transformer using Shifted Windows \(ai-contentlab.com\)](#)
- [Swin Transformer: Windows of Attention | by Subhash Nerella | Medium](#)
- [Swin Transformer - Paper Explained \(youtube.com\)](#)
- [Swin Transformer: Hierarchical Vision Transformer using Shifted Windows \(youtube.com\)](#)
- [The Question about the mask of window attention · Issue #38 · microsoft/Swin-Transformer \(github.com\)](#)
- [EfficientML.ai Lecture 14 - Vision Transformer \(MIT 6.5940, Fall 2023\) \(youtube.com\)](#)
- [Building Swin Transformer from Scratch using PyTorch: Hierarchical Vision Transformer using Shifted Windows | by Mishra | The Deep Hub | Medium](#)

# Implementation

- Instead of doing attention computation over all tokens, window attention does attention computation within fixed-size (e.g., 7x7) local windows.
- The number of tokens in each window is fixed. Thus, the computational complexity is linear.
- Gradually downsample the feature map size.
- Cross window information exchange
- SViT requires large data and compute expensive




# Train from transfer learning vs scratch

svit\_epoch\_0\_acc\_0.0000.pt  
svit\_epoch\_1\_acc\_0.8555.pt  
svit\_epoch\_2\_acc\_0.8926.pt  
svit\_epoch\_3\_acc\_0.9552.pt  
svit\_epoch\_8\_acc\_0.9796.pt  
svit\_epoch\_13\_acc\_0.9812.pt

Accuracy:  
0.98222232553694,  
Precision:  
0.98222232553694,  
Recall:  
0.98222232553694,  
F1 score:  
0.98222232553694

with pretrain weight: err 0/70  
Epochs: 30



ace\_of\_clubs\_46.jpg




four\_of\_clubs\_18.jpg



six\_of\_hearts\_29.jpg




three\_of\_spades\_42.jpg




eight\_of\_clubs\_21.jpg



joker\_68.jpg



ten\_of\_clubs\_45.jpg



two\_of\_spades\_10.jpg




eight\_of\_clubs\_59.jpg




nine\_of\_diamonds\_17.jpg



ten\_of\_spades\_57.jpg



five\_of\_clubs\_2.jpg



nine\_of\_diamonds\_34.jpg



three\_of\_clubs\_56.jpg

Accuracy:  
0.6305799550480313,  
Precision:  
0.6352797018157111,  
Recall:  
0.6305799550480313,  
F1 score:  
0.6224562724431356

from scratch: err 14/70  
Epochs: 100

svit\_epoch\_0\_acc\_0.0000.pt  
svit\_epoch\_1\_acc\_0.0619.pt  
svit\_epoch\_2\_acc\_0.1051.pt  
svit\_epoch\_3\_acc\_0.1566.pt  
svit\_epoch\_4\_acc\_0.1650.pt  
svit\_epoch\_5\_acc\_0.2104.pt  
svit\_epoch\_6\_acc\_0.2655.pt  
svit\_epoch\_7\_acc\_0.3060.pt  
svit\_epoch\_9\_acc\_0.3334.pt  
svit\_epoch\_10\_acc\_0.3558.pt  
svit\_epoch\_12\_acc\_0.4412.pt  
svit\_epoch\_13\_acc\_0.4776.pt  
svit\_epoch\_15\_acc\_0.5202.pt  
svit\_epoch\_20\_acc\_0.5293.pt  
svit\_epoch\_26\_acc\_0.5761.pt  
svit\_epoch\_28\_acc\_0.5765.pt  
svit\_epoch\_37\_acc\_0.6140.pt  
svit\_epoch\_43\_acc\_0.6168.pt  
svit\_epoch\_61\_acc\_0.6422.pt  
svit\_epoch\_65\_acc\_0.6617.pt  
svit\_epoch\_68\_acc\_0.6618.pt